Conditional random fields

upper row — generative model: model the prior probabilities of things

we're not asking how likely is it in general that i'd get a certain label. we can only answer the question: given these inputs, what's the probability of that label. we can't even in that case, answer the question given a label, what are the inputs?

— — more complicated in the sense that we have a whole joint distribution

lower row — discriminative world: model the conditional probabilities

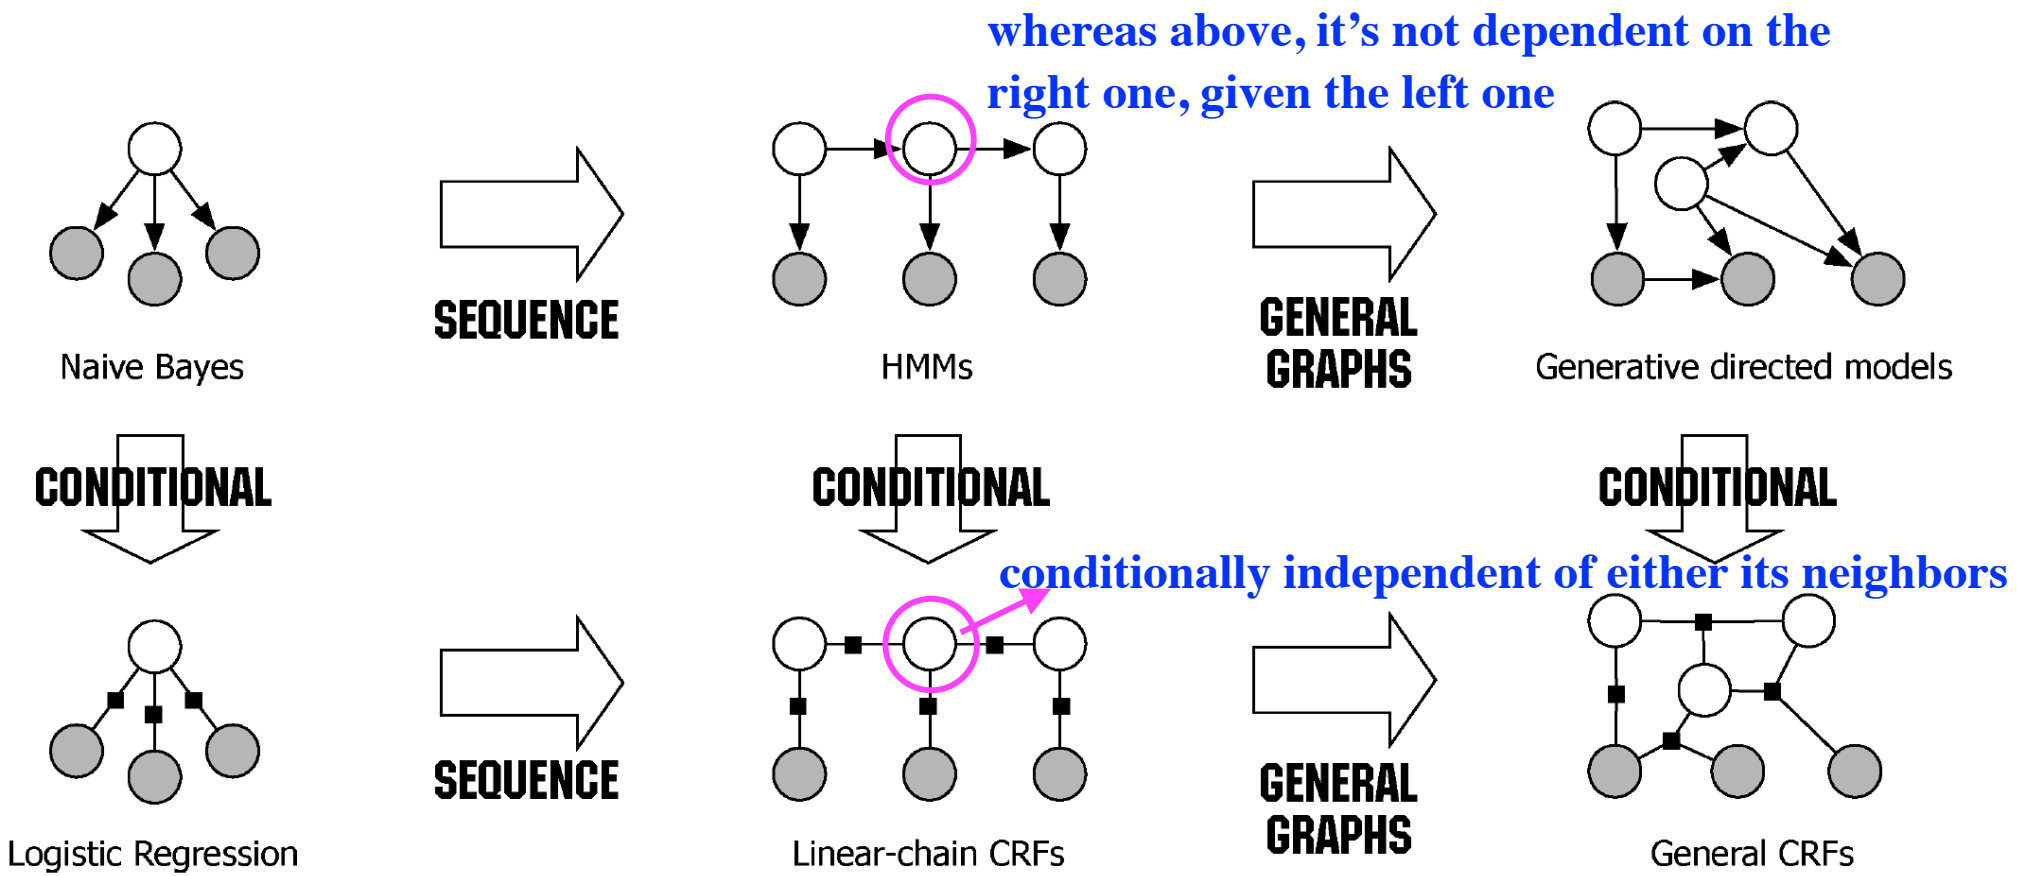

Fig. 2.4 Diagram of the relationship between naive Bayes, logistic regression, HMMs, linear-chain CRFs, generative models, and general CRFs.

Charles Sutton. An Introduction to Conditional Random Fields. Foundations and Trends in Machine Learning, 4(4):267–373, 2012.

CRF vs. Bayesian network

- discriminative
  - in contrast with generative models, CRF models only the conditional distribution $p(y|x)$
- undirected
- potentially cyclic

linear-chain CRF vs. HMM

- each label may be correlated with all observations
- feature functions may be arbitrary (differentiable)

HMM, restated

$$
\begin{aligned}
p(\mathbf{y}, \mathbf{x}) &= \frac{1}{Z} \prod_{t=1}^{T} \exp \left\{ \sum_{i,j \in S} \theta_{ij} \mathbf{1}_{\{y_t = i\}} \mathbf{1}_{\{y_{t-1} = j\}} + \sum_{i \in S} \sum_{o \in O} \mu_{oi} \mathbf{1}_{\{y_t = i\}} \mathbf{1}_{\{x_t = o\}} \right\} \\
&= \frac{1}{Z} \prod_{t=1}^{T} \exp \left\{ \sum_{i,j \in S} \theta_{ij} f_{ij}(y_t, y_{t-1}, x_t) + \sum_{i \in S} \sum_{o \in O} \mu_{oi} f_{io}(y_t, y_{t-1}, x_t) \right\} \\
&= \frac{1}{Z} \prod_{t=1}^{T} \exp \left\{ \sum_{k=1}^{K} \theta_k f_k(y_t, y_{t-1}, x_t) \right\}
\end{aligned}
$$

Linear-chain CRF

$$p(\mathbf{y}|\mathbf{x}) = \frac{1}{Z(\mathbf{x})} \prod_{t=1}^{T} \exp\left\{ \sum_{k=1}^{K} \theta_k f_k(y_t, y_{t-1}, \mathbf{x}_t) \right\}$$

Tasks

- can perform "decoding" - finding most likely state sequence via Viterbi
- can perform "inference" - finding best parameters θ_k via gradient descent

Application

feature functions can be:

- manually engineered features
    - e.g. "word ends in *-ing*"
- neural networks
    - e.g. the popular BiLSTM + CRF